

WEEK 1 PROGRESS UPDATE

Initial Exploration of Audio-Based Intent Understanding

Research Progress Report

Prepared by	Y.Unmesha Kruthika	Date	August 7, 2026
Reporting Period	Week 1	Status	Ongoing — Exploratory Phase

1. Introduction

The objective of this project is to explore methods for identifying the actual intention of a speaker from an audio input.

In human communication, the spoken words do not always represent the complete meaning intended by the speaker. The intention behind a sentence can depend on several factors such as context, tone, emotions, and the situation in which the statement is made.

Illustrative Example

Statement: "The room is cold."

The literal meaning indicates that the temperature is low. However, the actual intention may be:

- Requesting someone to switch off the air conditioner.
- Requesting a blanket.
- Asking someone to close a window.

Therefore, understanding the intention of a speaker requires going beyond simple speech recognition or text analysis.

The initial exploration focused on understanding existing approaches in:

- Spoken Language Understanding (SLU)
- Intent Detection
- Natural Language Understanding (NLU)
- Speech Representation Learning
- Multimodal approaches combining audio and text

2. Problem Understanding and Challenges

2.1 Indirect Intention Understanding

A major challenge in intention detection is that the literal meaning of a sentence may differ from the actual purpose of communication.

Humans naturally understand implied meanings using:

- Previous conversation
- Situational knowledge
- Tone of the speaker
- Social understanding

However, teaching machines to understand these hidden intentions remains challenging.

Spoken Sentence	Possible Actual Intention
"The bin is full."	Request to empty the bin
"It is dark here."	Request to switch on the light
"I am hungry."	Request for food

This problem requires systems that can understand the speaker's goal rather than only the words spoken.

2.2 Multilingual Speech Understanding

Another major challenge identified is handling multiple languages. Especially in countries like India, challenges include:

- Multiple languages and dialects
- Different meanings of words across regions
- Code-mixed speech (mixing multiple languages)
- Difficulty in preserving original meaning during translation

A speech-to-text followed by translation approach may sometimes change the original intention of the speaker. Therefore, multilingual intent understanding requires models capable of handling different languages while preserving the actual meaning.

3. Literature Study: Spoken Language Understanding (SLU)

A survey paper titled [“A Survey on Spoken Language Understanding: Recent Advancements and New Frontiers”](#) was studied to understand existing approaches and the evolution of intent detection systems. SLU is a field that focuses on understanding spoken language and extracting meaningful information from speech. The major components of SLU include:

3.1 Intent Detection

Intent detection identifies the goal or purpose behind a user's statement.

Field	Example
Input	“Book a flight ticket to Delhi.”
Intent	Flight Booking

The objective is to classify the user's intention rather than only identifying the words.

3.2 Slot Filling

Slot filling extracts important information related to the detected intent.

Sentence: “Book a flight from Hyderabad to Delhi tomorrow.”

Intent: Flight Booking

Slots:

- Source: Hyderabad
- Destination: Delhi
- Date: Tomorrow

Intent detection and slot filling together help systems understand user requests more effectively.

4. Evolution of Intent Detection Approaches

The literature survey helped understand the evolution of approaches used for intent detection.

4.1 Rule-Based Systems

Earlier systems relied on manually created rules.

Example: "If the sentence contains the word 'book', classify it as booking intent."

Limitations:

- Difficult to handle variations in language
- Required manual rule creation
- Poor adaptability

4.2 Traditional Machine Learning Approaches

Machine learning approaches improved automation by learning patterns from data. Common techniques included:

- Support Vector Machines (SVM)
- Naive Bayes

However, these approaches required manually extracted features.

4.3 Deep Learning Approaches

Recurrent Neural Networks (RNN)

RNNs were introduced for sequential data processing.

- Advantage: Can process sequences of words
- Limitation: Difficulty in handling long-term dependencies

Long Short-Term Memory Networks (LSTM)

LSTM improved RNNs by introducing memory mechanisms.

- Better handling of long-range dependencies
- Ability to remember important information

Gated Recurrent Units (GRU)

GRU is another improved version of RNN.

- Simpler architecture compared to LSTM
- Faster training while maintaining good performance

Convolutional Neural Networks (CNN)

CNN-based approaches were also applied for text classification.

- Advantage: Effective in identifying local patterns
- Limitation: Less effective in capturing long-range relationships compared to Transformer-based models

5. Transformer-Based Models for Intent Understanding

Recent advancements in Natural Language Processing have been dominated by Transformer-based models. Transformers introduced the attention mechanism, allowing models to understand relationships between different parts of a sentence.

Unlike RNN, LSTM, and GRU architectures, which process a sentence one word at a time, attention relates every word to every other word simultaneously — useful since the key information revealing intent may appear anywhere in the sentence.

Advantages:

- Better contextual understanding
- Ability to capture long-range dependencies
- Efficient processing of large datasets

BERT

Bidirectional Encoder Representations from Transformers.

BERT is called “bidirectional” because it learns a word's meaning from the words both before and after it, rather than reading only left to right, after first pretraining on large amounts of general text.

- Pretrained on large text datasets
- Learns language patterns
- Can be fine-tuned for tasks such as intent classification

RoBERTa

An improved version of BERT with better training strategies.


RoBERTa keeps BERT's architecture but refines the pretraining process — longer training on more data with a varied masking pattern — to produce more robust representations.

XLNet-RoBERTa

A multilingual Transformer model. Potentially useful for multilingual intent understanding because it supports multiple languages.

It is pretrained across many languages at once, learning a shared space where similar meanings sit close together regardless of language — directly relevant to the multilingual challenge in Section 2.2.

General Workflow



From the literature study, Transformer-based models are currently among the strongest approaches for text-based intent understanding.

6. Initial Exploration of Audio Understanding

Since the project focuses on audio input, understanding how AI processes speech was explored. A computer cannot directly understand audio — the speech signal must first be represented mathematically.

A raw audio recording is essentially a long sequence of numbers describing air-pressure variation over time — too long and too low-level for a model to learn from directly, so it must first be converted into a more compact representation. Three broad ways of doing this were studied.

6.1 Spectrogram-Based Representation

A spectrogram converts audio into a visual representation showing:

- Frequency variations
- Time information
- Signal intensity

In effect, a spectrogram turns a one-dimensional audio waveform into a two-dimensional, image-like representation, which lets it be processed using CNNs and other image-recognition networks.

6.2 MFCC Features

Mel Frequency Cepstral Coefficients (MFCCs) are traditional mathematical features extracted from audio. They represent speech characteristics based on human auditory perception. Used in:

- Speech recognition

- Speaker identification
- Audio classification


MFCCs are computed using a fixed, hand-designed procedure rather than being learned, making them compact and inexpensive — a useful lightweight baseline for this project.

6.3 Modern Speech Representation Models

wav2vec 2.0

wav2vec 2.0 is a self-supervised speech representation model. Instead of manually defining audio features, it learns useful speech representations directly from raw audio.

During pretraining, part of the audio signal is masked and the model learns to identify it from context, similar in spirit to how BERT masks words. This may preserve tone- and prosody-related information that is lost once speech is converted to plain text.




It can be fine-tuned for tasks such as:

- Speech recognition
- Emotion recognition
- Speech classification

7. Possible Multimodal Approach

Based on the initial exploration, combining multiple information sources may improve intention understanding.

Text alone tells us what was said, but the same words can carry different intentions depending on how they are spoken. Audio carries this additional layer through tone, stress, and pace. A possible approach:



Both parallel paths are combined at a fusion layer before a final prediction, so the outcome is informed by both what was said and how it was said.

Role of Each Component

Text Processing

Provides:

- Words
- Sentence meaning
- Linguistic information

Possible models:

- BERT
- RoBERTa

- XLM-RoBERTa

Audio Processing

Provides:

- Tone information
- Speaking style
- Acoustic patterns

Possible models:

- wav2vec 2.0
- Other speech representation models

Fusion Layer

The fusion layer combines the text-based and audio-based representations into a single representation before the final intent prediction. Common strategies include combining them early, combining them late after separate predictions, or learning attention weights that decide how much to rely on each path. Identifying the most suitable strategy is part of the next phase.

Context Information

Conversation history and surrounding information can help understand indirect intentions, and its exact role in the pipeline is part of the future exploration planned in Section 8.

8. Current Understanding and Future Exploration

Current Understanding

The initial study covered:

- Fundamentals of NLP, NLU, and SLU
- Intent detection and slot filling
- Evolution of intent detection approaches
- Transformer-based models for text understanding
- Audio representation techniques
- Modern speech models such as wav2vec 2.0
- Potential of multimodal approaches for intent understanding

Future Exploration Plan

The next phase will focus on:

1. Studying existing audio-based intent recognition systems.
2. Exploring research papers related to:
 - Audio + Text intent detection
 - Multimodal SLU
 - Multilingual intent understanding
 - Indirect intention detection
3. Identifying available datasets.
4. Understanding limitations of current systems.
5. Exploring possible implementation approaches.

Conclusion

The initial literature exploration indicates that intent understanding from audio is a challenging problem requiring multiple sources of information.

Text-based Transformer models provide strong language understanding capabilities, while speech representation models provide additional acoustic information.

A multimodal approach combining audio, text, and context may provide better understanding of the speaker's actual intention, especially in cases involving indirect meaning and multilingual communication.